

CV / VLMs

Unit 5: State-of-the-Art Object
Detection Techniques




5.2.1

Anchor-Free Object Detection


CenterNet and FCOS transformer
models

What is an Anchor Box? (recap)

- An anchor box is responsible for predicting an object class
- To effectively detect objects, a model has to try all possible anchor boxes for every grid cell



Anchor Box: a vehicle would be 1:1 (square) when looking from the front or rear, but 2:1 (rectangular) when viewed from the side.




Two anchor boxes (ships and airplane) are used to make two predictions per grid in the image.

Anchor-Free vs. Anchor-Based OD

While anchor-free object detection methods have equivalent accuracy to anchor-based methods, they offer several unique advantages:

1. **Simplified anchor identification:** Anchor-free methods eliminate the need to identify suitable anchors, which can be a complex problem due to factors like object physical properties, perspectives, and hyperparameters.
2. **Reduced computational complexity:** Anchor-based methods require more anchor boxes to improve accuracy, leading to more complex architectures and calculations. Anchor-free methods avoid this complexity.
3. **Improved generalizability:** Anchor-free object detection is more generalizable and can be extended to other tasks like keypoint detection and 3D object detection.

Speed vs Accuracy Plots




CenterNet latency vs mAP

Adapted from [Paper: Object as Points](#)

Introduction to CenterNet

In CenterNet, an object is represented by the center point (key-points) of its bounding box, which is crucial for its localization.



CenterNet: Objects as points.

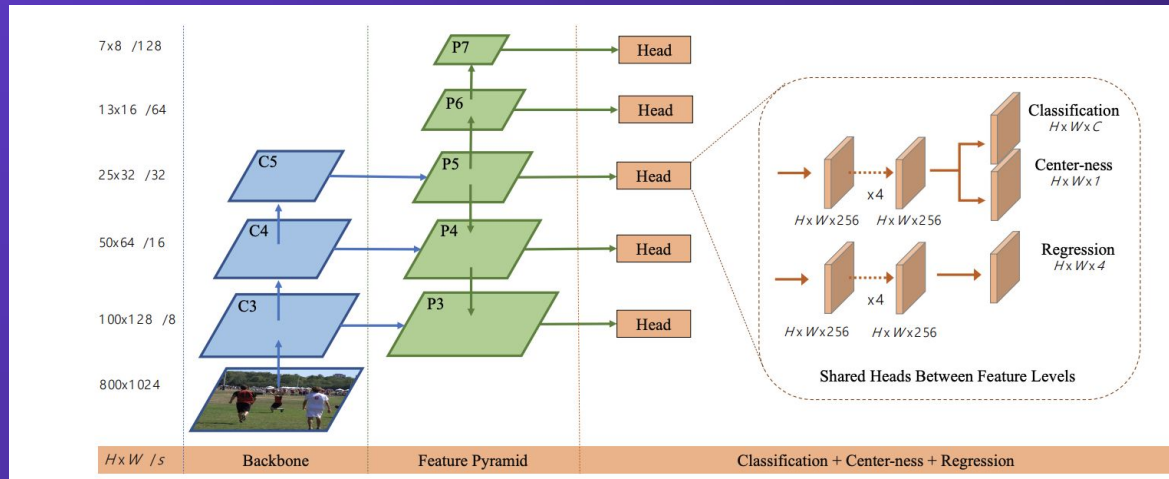
- **Keypoint heatmap:**
A set of heatmaps that would predict the likelihood of a pixel being the keypoint.
- **Local Offset:**
Offset to improve key-points precision
- **Bounding box size:**
Predicting width and height of the bounding box



CenterNet Components

Fully Convolutional One-Stage Object Detection (FCOS) Introduction

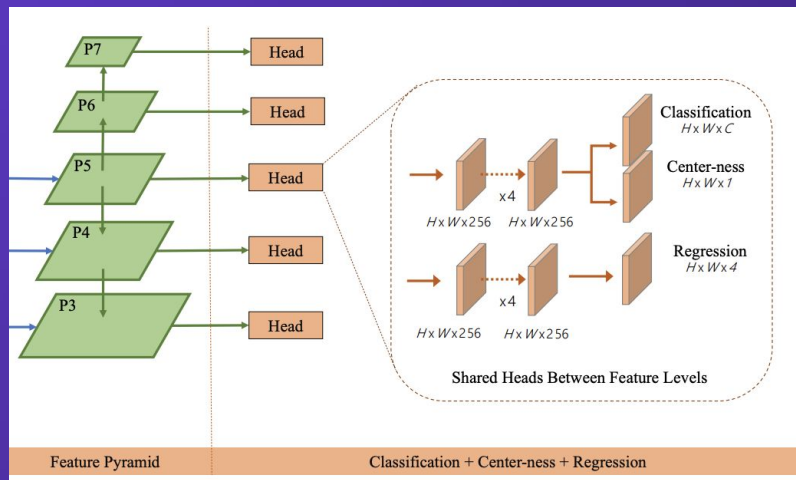
- FCOS (Fully Convolutional One-Stage Object Detection) uses a pixel-wise prediction for object detection.
- It utilizes FCN (Fully Convolutional Networks for semantic segmentations)
- Consist of 3 stages
 - Backbone
 - Feature Pyramid
 - Head
 - Classification loss
 - Center-ness loss
 - Regression loss



(top) network architecture of FCOS

Fully Convolutional One-Stage Object Detection (FCOS) Introduction (cont.)

- Classification
 - Predicting object classifications per spatial location.
 - Focal loss - modified standard cross entropy criterion
- Center-ness
 - Center-ness close to the bounding box center.
 - BCE (binary cross-entropy error/log) loss.
- Regression
 - Distance from center to edges (left, top, right, bottom)
 - IoU loss- which measures the overlap between the predicted and ground truth bounding boxes.



(top) network architecture of FCOS

