

# Building your First ASR / NLP Model




# Intro to NLP

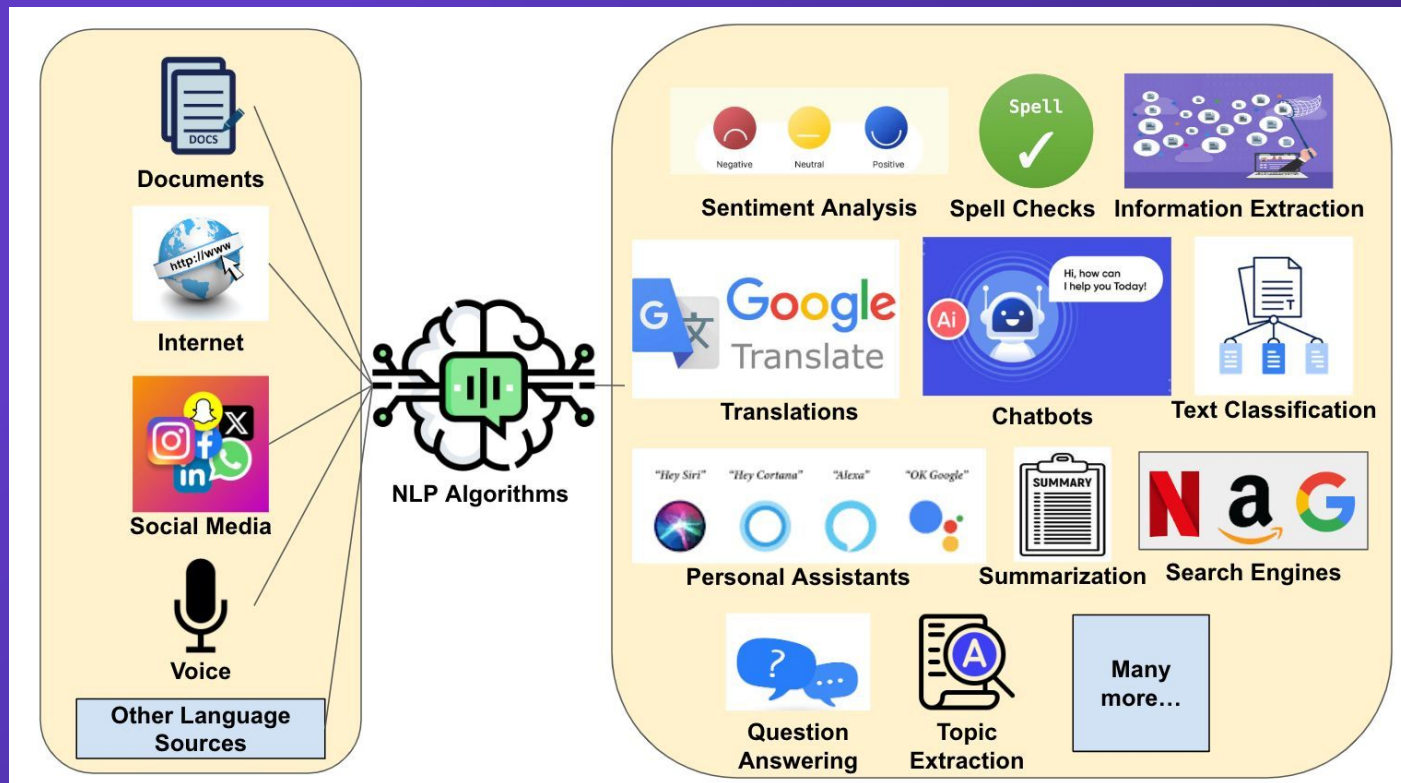
# What is NLP?

*Natural language processing (NLP) is a field at the intersection of computer science, artificial intelligence, and linguistics.*

*It concerns building systems that can process and understand human language.*




# Motivation behind NLP



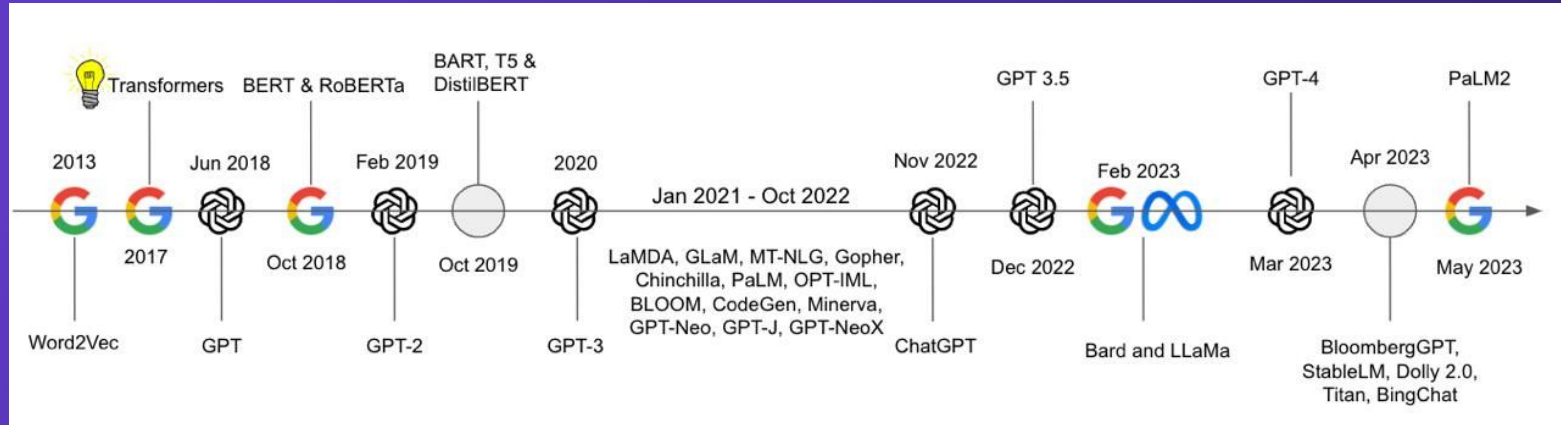
# Introduction to Question Answering

- Question Answering models can retrieve the answer to a question from a given text
- Useful for searching for an answer in a document



# What are Pre-trained Models

- Pre-trained models are massive deep learning models trained on enormous text datasets
- They have learned to represent language effectively




# Hands-on activity

# Intro to ASR



# What is Automatic Speech Recognition (ASR)?

- The process of converting spoken language into written text
- Powered by sophisticated computer algorithms
- Also known as speech-to-text or computer speech recognition



# Examples of ASR



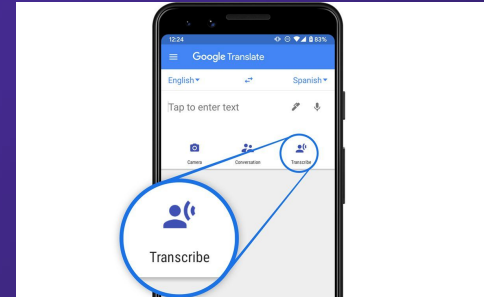
Virtual Assistants



Vehicle Speech Recognition



Accessibility for Hearing Impaired




Language Translation

# How we recognize speech?




# Time vs Frequency Domain

- ASR systems rely heavily on transforming speech signals from time domain to frequency domain to reveal the underlying frequencies that make up the speech
- Time domain: Speech is represented as a waveform where the amplitude (intensity) of the signal is plotted over time
- Discrete Fourier Transform (DFT) is used to decompose a time-domain signal into individual frequencies and amplitudes
- The output of the DFT tells us how much energy is present at each frequency in the original speech signal




# Spectrograms

- Visual representation of the spectrum of frequencies in a sound as they vary with time
- Shows how different frequencies appear, disappear, or change intensity over time in an audio signal



# Mel-Spectrograms

- Type of spectrogram where the frequency scale is converted to the Mel scale
- Mel scale more closely approximates human auditory system's response than the linear frequency scale; making it more effective for audio-related tasks in human speech and music



# Hands-on activity